

Mining Association Rules in Large Databases

Association rules

- Given a set of transactions D , find rules that will predict the occurrence of an item (or a set of items) based on the occurrences of other items in the transaction

Market-Basket transactions

<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Examples of association rules

$\{\text{Diaper}\} \rightarrow \{\text{Beer}\}$,
 $\{\text{Milk, Bread}\} \rightarrow \{\text{Diaper, Coke}\}$,
 $\{\text{Beer, Bread}\} \rightarrow \{\text{Milk}\}$,

An even simpler concept: frequent itemsets

- Given a set of transactions \mathbf{D} , find combination of items that occur frequently

Market-Basket transactions

<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Examples of frequent itemsets

{Diaper, Beer},
{Milk, Bread}
{Beer, Bread, Milk},

Lecture outline

- **Task 1:** Methods for finding all frequent itemsets efficiently
- **Task 2:** Methods for finding association rules efficiently

Definition: Frequent Itemset

- **Itemset**
 - A set of one or more items
 - E.g.: {Milk, Bread, Diaper}
 - **k-itemset**
 - An itemset that contains **k** items
- **Support count (σ)**
 - Frequency of occurrence of an itemset (number of transactions it appears)
 - E.g. $\sigma(\{\text{Milk, Bread, Diaper}\}) = 2$
- **Support**
 - Fraction of the transactions in which an itemset appears
 - E.g. $s(\{\text{Milk, Bread, Diaper}\}) = 2/5$
- **Frequent Itemset**
 - An itemset whose support is greater than or equal to a ***minsup*** threshold

<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke


Why do we want to find frequent itemsets?

- Find all combinations of items that occur together
- They might be interesting (e.g., in placement of items in a store ☺)
- Frequent itemsets are only positive combinations (we do not report combinations that do not occur frequently together)
- Frequent itemsets aims at providing a summary for the data

Finding frequent sets

- **Task:** Given a transaction database D and a **minsup** threshold find all frequent itemsets and the frequency of each set in this collection
- **Stated differently:** Count the number of times combinations of attributes occur in the data. If the count of a combination is above **minsup** report it.
- **Recall:** The input is a transaction database D where every transaction consists of a subset of items from some universe I

How many itemsets are there?




Given d items, there
are 2^d possible
itemsets

When is the task sensible and feasible?

- If **minsup = 0**, then all subsets of I will be frequent and thus the size of the collection will be very large
- This summary is very large (maybe larger than the original input) and thus not interesting
- The task of finding all frequent sets is interesting typically only for relatively large values of **minsup**

A simple algorithm for finding all frequent itemsets ??



Brute-force algorithm for finding all frequent itemsets?

- Generate all possible itemsets (lattice of itemsets)
 - Start with 1-itemsets, 2-itemsets,...,d-itemsets
- Compute the frequency of each itemset from the data
 - Count in how many transactions each itemset occurs
- If the support of an itemset is above **minsup** report it as a frequent itemset

Brute-force approach for finding all frequent itemsets

- Complexity?
 - Match every candidate against each transaction
 - For **M** candidates and **N** transactions, the complexity is $\sim O(NM^w)$ => **Expensive since $M = 2^d$!!!**

Speeding-up the brute-force algorithm

- Reduce the **number of candidates (M)**
 - Complete search: $M=2^d$
 - Use pruning techniques to reduce M
- Reduce the **number of transactions (N)**
 - Reduce size of N as the size of itemset increases
 - Use vertical-partitioning of the data to apply the mining algorithms
- Reduce the **number of comparisons (NM)**
 - Use efficient data structures to store the candidates or transactions
 - No need to match every candidate against every transaction

Reduce the number of candidates

- **Apriori principle (Main observation):**
 - If an itemset is frequent, then all of its subsets must also be frequent
- Apriori principle holds due to the following property of the support measure:

$$\forall X, Y : (X \subseteq Y) \Rightarrow s(X) \geq s(Y)$$

- The support of an itemset **never exceeds** the support of its subsets
- This is known as the ***anti-monotone*** property of support

Example


<i>TID</i>	<i>Items</i>
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

$$s(\text{Bread}) > s(\text{Bread, Beer})$$

$$s(\text{Milk}) > s(\text{Bread, Milk})$$

$$s(\text{Diaper, Beer}) > s(\text{Diaper, Beer, Coke})$$

Illustrating the Apriori principle



Illustrating the Apriori principle

Item	Count
Bread	4
Coke	2
Milk	4
Beer	3
Diaper	4
Eggs	1

Items (1-itemsets)

Itemset	Count
{Bread,Milk}	3
{Bread,Beer}	2
{Bread,Diaper}	3
{Milk,Beer}	2
{Milk,Diaper}	3
{Beer,Diaper}	3

Pairs (2-itemsets)

(No need to generate candidates involving Coke or Eggs)

minsup = 3/5

If every subset is considered,

$${}^6C_1 + {}^6C_2 + {}^6C_3 = 41$$

With support-based pruning,

$$6 + 6 + 1 = 13$$

Item set	Count
{Bread,Milk,Diaper}	3

Triplets (3-itemsets)

Exploiting the Apriori principle

1. Find **frequent 1-items** and put them to L_k ($k=1$)
2. Use L_k to generate a collection of *candidate* itemsets C_{k+1} with size ($k+1$)
3. Scan the database to find which itemsets in C_{k+1} are **frequent** and put them into L_{k+1}
4. If L_{k+1} is not empty
 - $k=k+1$
 - Goto step 2

The Apriori algorithm

C_k : Candidate itemsets of size k

L_k : frequent itemsets of size k

L_1 = {frequent 1-itemsets};

for ($k = 2; L_k \neq \emptyset; k++$)

C_{k+1} = GenerateCandidates(L_k)

for each transaction t in database **do**

increment count of candidates in C_{k+1} that are contained in t

endfor

L_{k+1} = candidates in C_{k+1} with support $\geq min_sup$

endfor

return $\cup_k L_k$;

GenerateCandidates

- Assume the items in L_k are listed in an order (e.g., alphabetical)
- **Step 1: self-joining L_k (IN SQL)**

insert into C_{k+1}


select $p.item_1, p.item_2, \dots, p.item_k, q.item_k$

from $L_k p, L_k q$

where $p.item_1=q.item_1, \dots, p.item_{k-1}=q.item_{k-1}, p.item_k < q.item_k$

Example of Candidates Generation

- $L_3 = \{abc, abd, acd, ace, bcd\}$
- ***Self-joining:*** $L_3 * L_3$
 - ***abcd*** from ***abc*** and ***abd***
 - ***acde*** from ***acd*** and ***ace***



GenerateCandidates

- Assume the items in L_k are listed in an order (e.g., alphabetical)
- **Step 1:** *self-joining L_k (IN SQL)*

insert into C_{k+1}

select $p.item_1, p.item_2, \dots, p.item_k, q.item_k$

from $L_k p, L_k q$

where $p.item_1=q.item_1, \dots, p.item_{k-1}=q.item_{k-1}, p.item_k < q.item_k$

- **Step 2:** *pruning*

forall *itemsets c in C_{k+1}* do

 forall *k-subsets s of c* do

 if (*s is not in L_k*) then delete *c* from C_{k+1}

Example of Candidates Generation

- $L_3 = \{abc, abd, acd, ace, bcd\}$


- ***Self-joining:*** $L_3 * L_3$

- ***abcd*** from ***abc*** and ***abd***
 - ***acde*** from ***acd*** and ***ace***

- ***Pruning:***

- ***acde*** is removed because ***ade*** is not in L_3

- $C_4 = \{abcd\}$



The Apriori algorithm

C_k : Candidate itemsets of size k

L_k : frequent itemsets of size k

L_1 = {frequent items};

for ($k = 1$; $L_k \neq \emptyset$; $k++$)

C_{k+1} = GenerateCandidates(L_k)

for each transaction t in database **do**

increment count of candidates in C_{k+1} that are contained in t

endfor

L_{k+1} = candidates in C_{k+1} with support $\geq min_sup$

endfor


return $\cup_k L_k$;

How to Count Supports of Candidates?

- Naive algorithm?
- Method:
 - Candidate itemsets are stored in a *hash-tree*
 - *Leaf node* of hash-tree contains a list of itemsets and counts
 - *Interior node* contains a hash table
 - *Subset function*: finds all the candidates contained in a transaction


Example of the hash-tree for C_3

Hash function: mod 3




Example of the hash-tree for C_3

Hash function: mod 3




Hash on 3rd item



Example of the hash-tree for C_3

Hash function: mod 3



The subset function finds all the candidates contained in a transaction:

- At the root level it hashes on all items in the transaction
- At level i it hashes on all items in the transaction that come after item the i -th item

Discussion of the Apriori algorithm

- Much faster than the Brute-force algorithm
 - It avoids checking all elements in the lattice
- The running time is in the worst case $O(2^d)$
 - Pruning really prunes in practice
- It makes multiple passes over the dataset
 - One pass for every level k
- Multiple passes over the dataset is inefficient when we have thousands of candidates and millions of transactions


Making a single pass over the data: the AprioriTid algorithm

- The database is **not** used for counting support after the 1st pass!
- Instead information in data structure C'_k is used for counting support in every step
 - $C'_k = \{<\text{TID}, \{X_k\}> \mid X_k \text{ is a potentially frequent } k\text{-itemset in transaction with id=TID}\}$
 - C'_1 : corresponds to the original database (every item i is replaced by itemset $\{i\}$)
 - The member C'_k corresponding to transaction t is $<t.\text{TID}, \{c \in C_k \mid c \text{ is contained in } t\}>$

The AprioriTID algorithm

- $L_1 = \{\text{frequent 1-itemsets}\}$
- $C_1' = \text{database } D$
- **for** ($k=2, L_{k-1}' \neq \text{empty}; k++$)
 - $C_k = \text{GenerateCandidates}(L_{k-1})$
 - $C_k' = \{\}$
 - for** all entries $t \in C_{k-1}'$
 - $C_t = \{c \in C_k \mid t[c-c[k]] = 1 \text{ and } t[c-c[k-1]] = 1\}$
 - for** all $c \in C_t$ { $c.\text{count}++$ }
 - if** ($C_t \neq \{\}$)
 - append* C_t to C_k'
 - endif**
 - endfor**
 - $L_k = \{c \in C_k \mid c.\text{count} \geq \text{minsup}\}$
- **endfor**
- **return** $\bigcup_k L_k$

AprioriTid Example (minsup=2)



Discussion on the AprioriTID algorithm

- L_1 = {frequent 1-itemsets}
 - C_1' = database D
 - **for** ($k=2$, $L_{k-1}' \neq$ empty; $k++$)
 - C_k = GenerateCandidates(L_{k-1}')
 - $C_k' = \{\}$
 - for** all entries $t \in C_{k-1}'$
 - $C_t = \{c \in C_k \mid t[c-c[k]] = 1 \text{ and } t[c-c[k-1]] = 1\}$
 - for** all $c \in C_t$ { $c.\text{count}++$ }
 - if** ($C_t \neq \{\}$)
 - append* C_t to C_k'
 - endif**
 - endfor**
 - $L_k = \{c \in C_k \mid c.\text{count} \geq \text{minsup}\}$
 - endfor**
 - **return** $U_k L_k$
- One single pass over the data
 - C_k' is generated from C_{k-1}'
 - For small values of k , C_k' could be larger than the database!
 - For large values of k , C_k' can be very small

Apriori vs. AprioriTID

- *Apriori* makes multiple passes over the data while *AprioriTID* makes a single pass over the data
- *AprioriTID* needs to store additional data structures that may require more space than *Apriori*
- Both algorithms need to check all candidates' frequencies in every step

Lecture outline

- **Task 1:** Methods for finding all frequent itemsets efficiently
- **Task 2:** Methods for finding association rules efficiently

Definition: Association Rule

Let D be database of **transactions**

– e.g.:

Transaction ID	Items
2000	A, B, C
1000	A, C
4000	A, D
5000	B, E, F

- Let I be the set of items that appear in the database, e.g., $I = \{A, B, C, D, E, F\}$
- A **rule** is defined by $X \rightarrow Y$, where $X \subset I$, $Y \subset I$, and $X \cap Y = \emptyset$
 - e.g.: $\{B, C\} \rightarrow \{A\}$ is a rule

Definition: Association Rule

□ Association Rule

- An implication expression of the form $X \rightarrow Y$, where X and Y are non-overlapping itemsets
- Example:
 $\{Milk, Diaper\} \rightarrow \{Beer\}$

□ Rule Evaluation Metrics

- Support (s)
 - Fraction of transactions that contain both X and Y
- Confidence (c)
 - Measures how often items in Y appear in transactions that contain X

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke


Example:

$$\{Milk, Diaper\} \rightarrow Beer$$

$$s = \frac{\sigma(Milk, Diaper, Beer)}{|T|} = \frac{2}{5} = 0.4$$

$$c = \frac{\sigma(Milk, Diaper, Beer)}{\sigma(Milk, Diaper)} = \frac{2}{3} = 0.67$$

Rule Measures: Support and Confidence



Find all the rules $X \rightarrow Y$ with minimum confidence and support

- support, s , probability that a transaction contains $\{X \cup Y\}$
- confidence, c , *conditional probability* that a transaction having X also contains Y

TID	Items
100	A,B,C
200	A,C
300	A,D
400	B,E,F

Let minimum support 50%, and minimum confidence 50%, we have

- $A \rightarrow C$ (50%, 66.6%)
- $C \rightarrow A$ (50%, 100%)

Example

TID	date	items bought
100	10/10/99	{F,A,D,B}
200	15/10/99	{D,A,C,E,B}
300	19/10/99	{C,A,B,E}
400	20/10/99	{B,A,D}

What is the *support* and *confidence* of the rule: $\{B,D\} \rightarrow \{A\}$

- Support:
 - percentage of tuples that contain $\{A,B,D\}$ = 75%
- Confidence:
$$\frac{\text{number of tuples that contain } \{A,B,D\}}{\text{number of tuples that contain } \{B,D\}} = 100\%$$

Association-rule mining task


- Given a set of transactions **D**, the goal of association rule mining is to find **all** rules having
 - support $\geq \textit{minsup}$ threshold
 - confidence $\geq \textit{minconf}$ threshold

Brute-force algorithm for association-rule mining

- List all possible association rules
- Compute the support and confidence for each rule
- Prune rules that fail the *minsup* and *minconf* thresholds
- \Rightarrow Computationally prohibitive!

Computational Complexity

- Given d unique items in I :
 - Total number of itemsets = 2^d
 - Total number of possible association rules:



$$R = \sum_{k=1}^{d-1} \binom{d}{k} \times \sum_{j=1}^{d-k} \binom{d-k}{j}$$
$$= 3^d - 2^{d+1} + 1$$

Mining Association Rules

TID	Items
1	Bread, Milk
2	Bread, Diaper, Beer, Eggs
3	Milk, Diaper, Beer, Coke
4	Bread, Milk, Diaper, Beer
5	Bread, Milk, Diaper, Coke

Example of Rules:

$\{\text{Milk}, \text{Diaper}\} \rightarrow \{\text{Beer}\}$ ($s=0.4, c=0.67$)
 $\{\text{Milk}, \text{Beer}\} \rightarrow \{\text{Diaper}\}$ ($s=0.4, c=1.0$)
 $\{\text{Diaper}, \text{Beer}\} \rightarrow \{\text{Milk}\}$ ($s=0.4, c=0.67$)
 $\{\text{Beer}\} \rightarrow \{\text{Milk}, \text{Diaper}\}$ ($s=0.4, c=0.67$)
 $\{\text{Diaper}\} \rightarrow \{\text{Milk}, \text{Beer}\}$ ($s=0.4, c=0.5$)
 $\{\text{Milk}\} \rightarrow \{\text{Diaper}, \text{Beer}\}$ ($s=0.4, c=0.5$)

Observations:

- All the above rules are binary partitions of the same itemset:
 $\{\text{Milk}, \text{Diaper}, \text{Beer}\}$
- Rules originating from the same itemset have identical support but can have different confidence
- Thus, we may decouple the support and confidence requirements

Mining Association Rules

- Two-step approach:
 - **Frequent Itemset Generation**
 - Generate all itemsets whose support $\geq \text{minsup}$
 - **Rule Generation**
 - Generate high confidence rules from each frequent itemset, where each rule is a binary partition of a frequent itemset

Rule Generation – Naive algorithm

- Given a frequent itemset X , find all non-empty subsets $y \subset X$ such that $y \rightarrow X - y$ satisfies the minimum confidence requirement
 - If $\{A, B, C, D\}$ is a frequent itemset, candidate rules:

$ABC \rightarrow D,$	$ABD \rightarrow C,$	$ACD \rightarrow B,$	$BCD \rightarrow A,$
$A \rightarrow BCD,$	$B \rightarrow ACD,$	$C \rightarrow ABD,$	$D \rightarrow ABC$
$AB \rightarrow CD,$	$AC \rightarrow BD,$	$AD \rightarrow BC,$	$BC \rightarrow AD,$
$BD \rightarrow AC,$	$CD \rightarrow AB,$		
- If $|X| = k$, then there are $2^k - 2$ candidate association rules (ignoring $L \rightarrow \emptyset$ and $\emptyset \rightarrow L$)

Efficient rule generation


- How to efficiently generate rules from frequent itemsets?
 - In general, confidence does not have an anti-monotone property
 $c(ABC \rightarrow D)$ can be larger or smaller than $c(AB \rightarrow D)$
 - *But confidence of rules generated from the same itemset has an anti-monotone property*
 - Example: $X = \{A, B, C, D\}$:
 $c(ABC \rightarrow D) \geq c(AB \rightarrow CD) \geq c(A \rightarrow BCD)$
 - Why?

Confidence is anti-monotone w.r.t. number of items on the RHS of the rule

Rule Generation for Apriori Algorithm

Lattice of rules

Low
Confidence
Rule



Apriori algorithm for rule generation

- Candidate rule is generated by merging two rules that share the same prefix in the rule consequent
- **join($CD \rightarrow AB, BD \rightarrow AC$)** would produce the candidate rule $D \rightarrow ABC$
- **Prune** rule $D \rightarrow ABC$ if there exists a subset (e.g., $AD \rightarrow BC$) that does not have high confidence

